

# UNIT - 2

## PROBLEM SOLVING AND REPRESENTATION OF KNOWLEDGE

By Vishvajit Bakrola

# OUTLINE

- ❑ Problem Solving Agents
- ❑ Searching for Solutions
- ❑ Search Techniques – Uninformed and Informed
- ❑ Heuristic Functions
- ❑ Local Search and Optimization
- ❑ Knowledge based Agents
- ❑ Logics



# PROBLEM SOLVING AGENTS

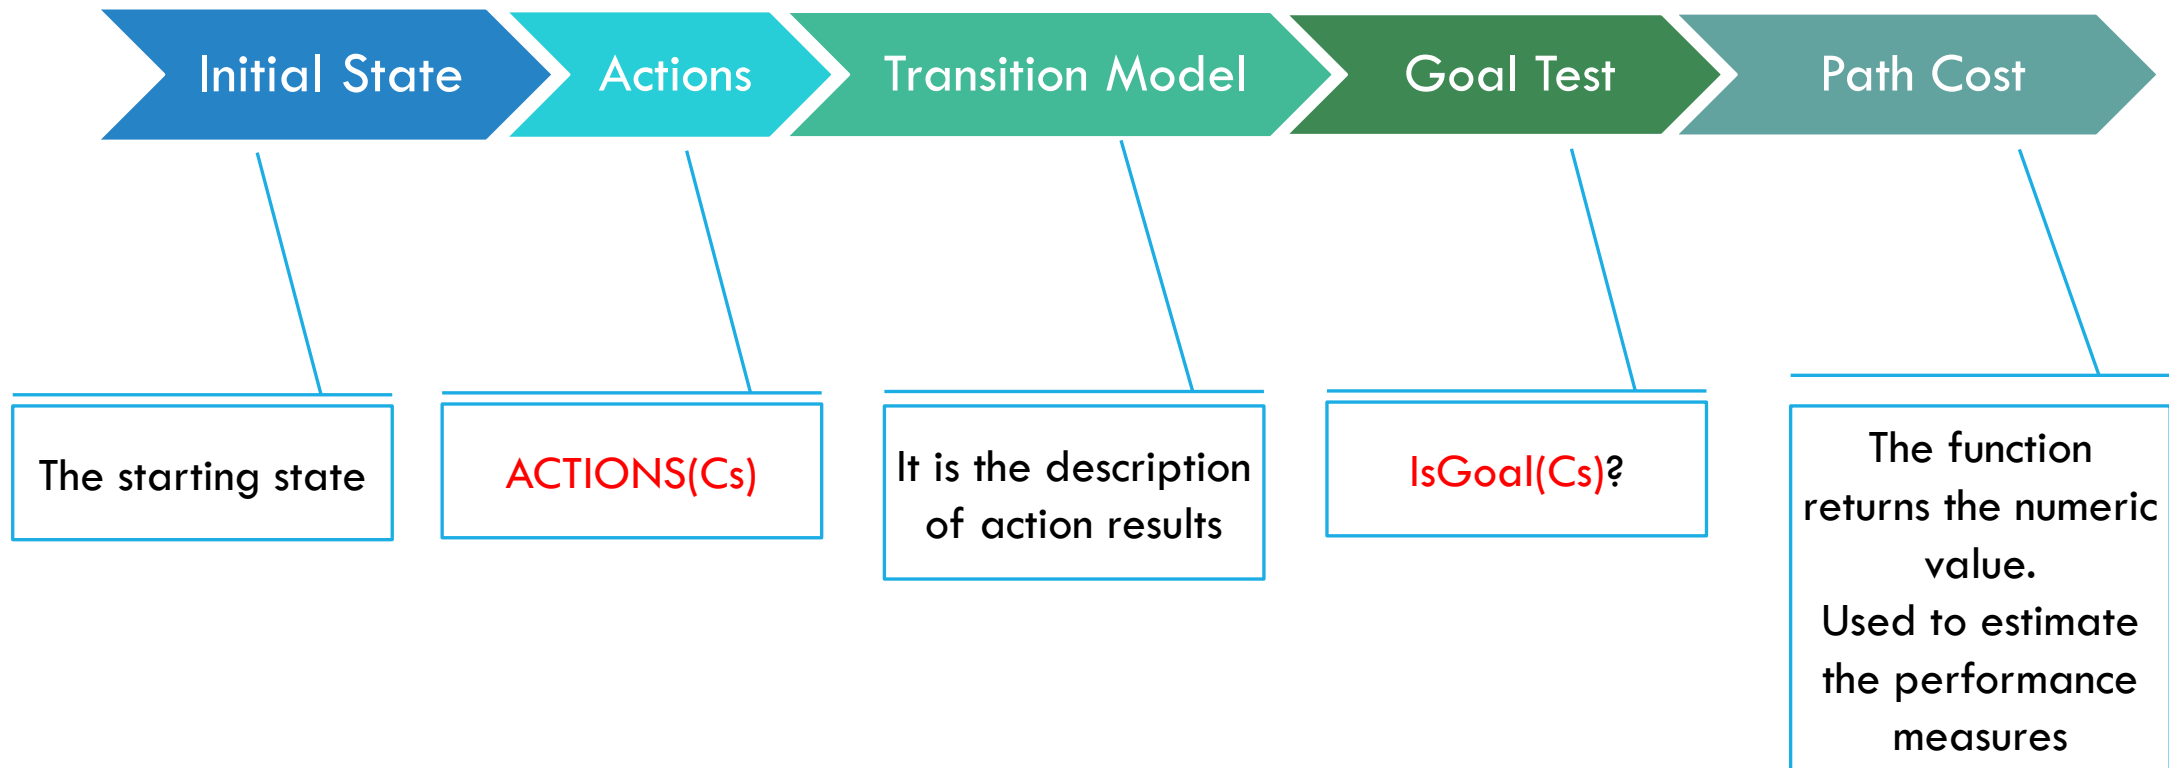
# PROBLEM SOLVING AGENTS

- They are **Goal based** or **Goal directed** agents.
- We know that, an intelligent agent are supposed to maximize their performance measures.
- **Goal formation** is the first step of any problem solving agent.
- A goal can be a state or a path (*i.e. sequential collection of states*).
- The process of looking for a sequence of actions that leads an agent to the goal is known as **Search**.
- Search algorithm takes a problem as an input and returns **solution** having the **sequence of actions**.

# WELL-DEFINED PROBLEMS AND SOLUTIONS

- A problem can be systematically defined using following five components:

**Optimal solution should have lowest path cost**



# REAL-WORLD EXAMPLE

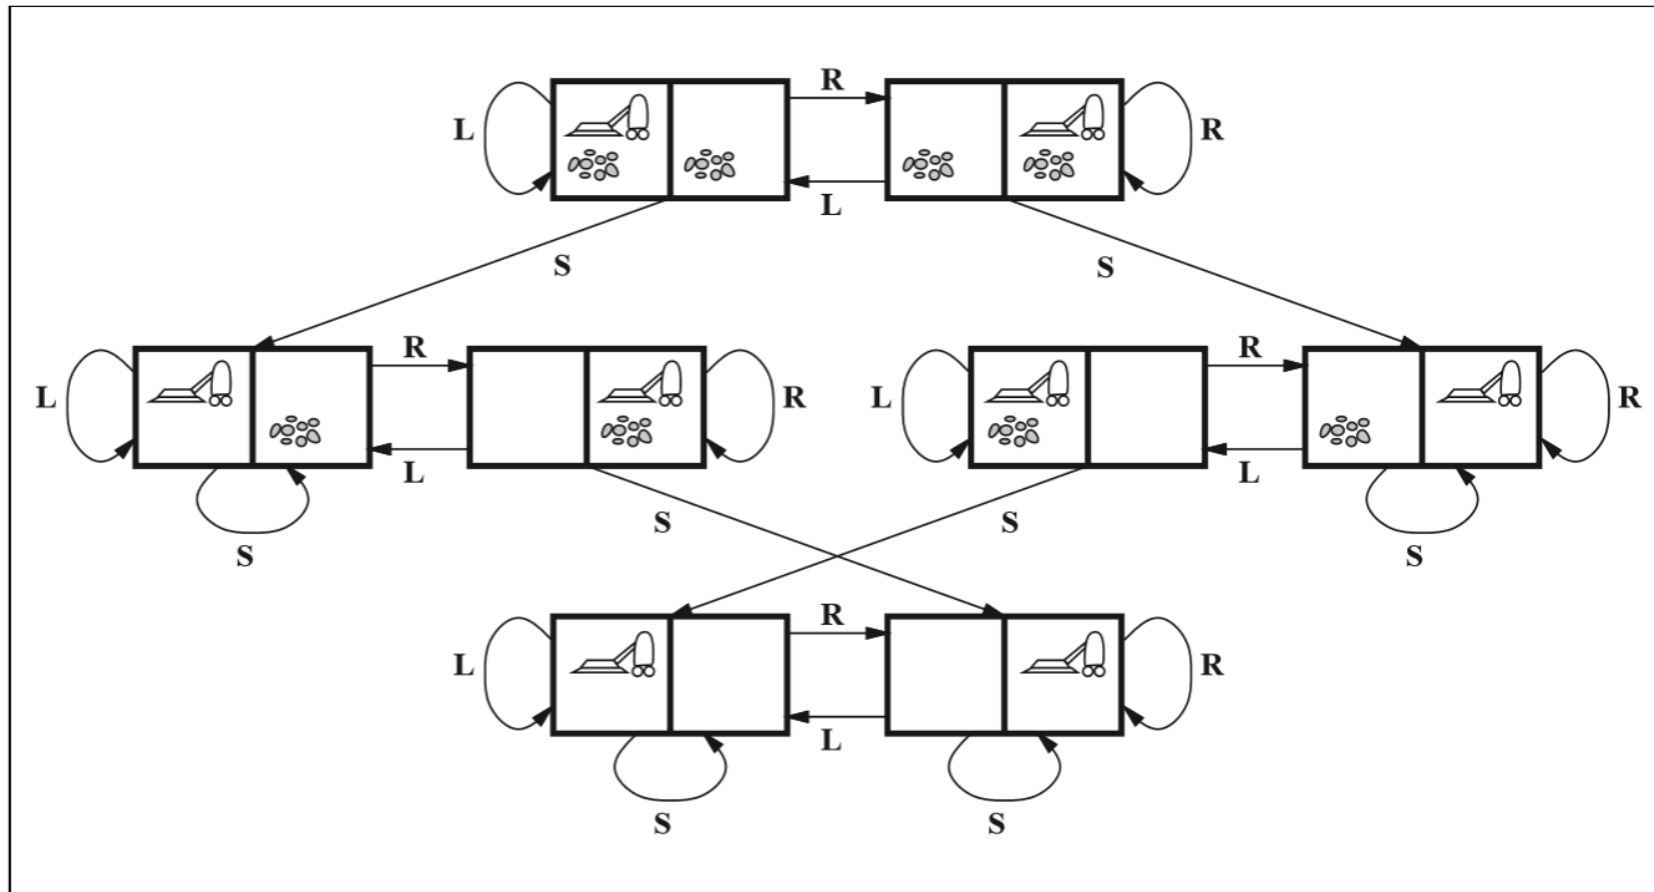


Figure 1: The state space of vacuum cleaner world

# PROBLEM FORMULATION

- **States:** The state is determined by **both the agent location and the dirt locations**. The agent is in one of two locations, each of which might or might not contain dirt. Thus, there are  $2 \times 2^2 = 8$  possible world states. A larger environment with  $n$  locations has  $n \times 2^n$  states.
- **Initial state:** Any state can be designated as the initial state.
- **Actions:** In this simple environment, each state has just three actions: Left, Right, and Suck. Larger environments might also include Up and Down.
- **Transition model:** The actions have their expected effects, except that moving *Left* in the leftmost square, moving *Right* in the rightmost square, and *Sucking* in a clean square have no effect.

# CONTINUE...

- **Goal test:** This checks whether all the squares are clean.
- **Path cost:** Each step costs 1, so the path cost is the number of steps in the path.



# THE 8-PUZZLE EXAMPLE

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

Figure 2: A typical instance of 8-puzzle problem

# EXERCISE

- State??
- Initial state??
- Actions??
- Transition model??
- Goal test??
- Path cost??

# CONTINUE...

- Other examples can be,
  - ✓ Sliding block puzzle
  - ✓ 8-queen problem
- Some real-world problems
  - ✓ Rout-finding problems
  - ✓ Touring problems
  - ✓ TSP
  - ✓ VLSI layout designing
  - ✓ Robot navigation
  - ✓ Assembly sequencing




**SEARCHING FOR THE SOLUTION**

# SEARCHING OF THE SOLUTION

- As we know, solution is a **sequence of actions**.
- Select the initial node (*if possibilities are many*)
- Then we apply actions to current node, with possible legal actions and generate new set of states.
- If there is no goal state from newly generate set of states, then the same process will be repeated until we reach goal state.
- In last, the search algorithm will return goal state/solution.

# CONTINUE...



Root Node/Current State  
Check **IsGoal()**  
Action = **Expand(0)**

Nodes/States at next level  
Check **IsGoal()**  
Action = **Expand(0)**

Nodes/States at next level  
Check **IsGoal()**  
If goal found, return goal  
Terminate the search

Figure 3: Search tree expansion

# INFRASTRUCTURE FOR SEARCH ALGORITHMS

- Search algorithms require a data structure to keep track of the search tree that is being constructed. For each node  $n$  of the tree, we have a structure that contains four components:
  1. **n.STATE**: the state in the state space to which the node corresponds.
  2. **n.PARENT**: the node in the search tree that generated this node.
  3. **n.ACTION**: the action that was applied to the parent to generate the node.
  4. **n.PATH-COST**: the cost, traditionally denoted by  $g(n)$ , of the path from the initial state to the node, as indicated by the parent pointers.

# MEASURING PROBLEM SOLVING PERFORMANCE

- We can evaluate an algorithm's performance in four ways:
  - 1. Completeness:** Is the algorithm guaranteed to find a solution when there is one?
  - 2. Optimality:** Does the strategy find the optimal solution?
  - 3. Time complexity:** How long does it take to find a solution?
  - 4. Space complexity:** How much memory is needed to perform the search?



# CONTINUE...

- Time and space complexity are always considered with respect to some measure of the problem difficulty.
- The typical measure is the size of the state space graph,  $|V| + |E|$ , where  $V$  is the set of vertices (nodes) of the graph and  $E$  is the set of edges (links).
- For these reasons, complexity is expressed in terms of three quantities:  $b$ , the **branching factor** or maximum number of successors of any node;  $d$ , the **depth** of the shallowest goal node (i.e., the number of steps along the path from the root); and  $m$ , the maximum length of any path in the state space.
- To assess the effectiveness of a search algorithm, we can consider just the **search cost**—which typically depends on the time complexity but can also include a term for memory usage—or we can use the **total cost**, which combines the search cost and the path cost of the solution found.



# UNINFORMED SEARCH STRATEGIES

# SOME IMPORTANT TERMS

- ❑ **Search space** → possible conditions and solutions.
- ❑ **Initial state** → state where the searching process started.
- ❑ **Goal state** → the ultimate aim of searching process.
- ❑ **Problem space** → “what to solve”
- ❑ **Searching strategy** → strategy for controlling the search.
- ❑ **Search tree** → tree representation of search space, showing possible solutions from initial state.

# SEARCHING STRATEGIES

- **Blind search** → traversing the search space until the goal nodes is found (might be doing exhaustive search).

- Techniques : **Breadth First Uniform Cost ,Depth first, Interactive Deepening search.**

- Guarantees solution.

- **Heuristic search** → search process takes place by traversing search space with applied rules (information).





- Techniques: **Greedy Best First Search, A\* Algorithm**

- There is no guarantee that solution is found.

# BREADTH FIRST SEARCH (BFS)

- **Strategy:** Search all the nodes expanded at given depth before any node at next level.
- **Concept :** First In First Out (FIFO) queue.
- **Complete ?:** Yes with finite  $b$  (branch).
- **Space :** Keep nodes in every memory
- **Optimal ?** = Yes (if cost = 1)







# BFS - EXAMPLE



# DEPTH FIRST SEARCH (DFS)

- **Strategy:** Search all the nodes expanded in deepest path.
- **Concept:** Last In First Out
- **Complete ?:** No
- **Complexity:**  $O(b^m)$
- **Space :**  $O(bm)$  –  $b$  ; *branching factor*,  $m$  ; *max. depth*
- **Optimality ? :** No

# DFS - EXAMPLE





# DEPTH-LIMITED SEARCH (DLS)

- DLS is DFS with a depth limit,  $L$ 
  - ✓ Nodes at depth limit are treated as leaves
  - ✓ Depth limits can come from domain knowledge
    - E.g.,  $L = \text{diameter}$  of state space
- Implemented same as DFS, but with depth limit
- Performance
  - ✓ Incomplete if  $L < d$
  - ✓ Nonoptimal if  $L > d$
  - ✓ Time and space complexity  $\leq$  DFS

# ITERATIVE DEEPENING DFS (ID-DFS)


- IDDFS performs DLS one or more times, each time with a larger depth limit,  $l$ 
  - ✓ Start with  $l = 0$ , perform DLS
  - ✓ If goal not found, repeat DLS with  $l = 1$ , and so on
- Performance
  - ✓ Combines benefits of BFS (completeness, optimality) with benefits of DFS (relatively low space complexity)
- Not as inefficient as it looks
  - ✓ Repeating low-depth levels: these levels do not have many nodes

# ID-DFS - EXAMPLE

Limit = 1




Limit = 3



# BIDIRECTIONAL SEARCH

- Perform two searches simultaneously
  - ✓ One search from the initial state to a goal state
  - ✓ Another search from a goal state to the initial state
  - ✓ Stop searching when either search reaches a state that is in the fringe of the other state (i.e., when they “meet in the middle”)
- Performance
  - ✓ Complete and optimal if both searches are BFS

# CONTINUE...



Order of visitation: 1, 2, 3, ...

Figure 4: Bidirectional Search Example

# CONTINUE...

- Reaching a node with a repeated state means that there are at least two paths to the same state
  - ✓ Can lead to infinite loops
- Keep a list of all nodes expanded so far (called the closed list)



# **INFORMED SEARCH AND EXPLORATION**

# INFORMED SEARCH AND EXPLORATION

- Review of Associated Functions
- Greedy Best-First Search
- A\* Search
- Heuristic Functions



# REVIEW OF ASSOCIATED FUNCTIONS

- $g(n)$  = cost from the initial state to the current state  $n$
- $h(n)$  = estimated cost of the cheapest path from node  $n$  to a goal node
- $f(n)$  = evaluation function to select a node for expansion (usually the lowest cost node)

# GREEDY BEST-FIRST SEARCH

- Greedy Best-First search tries to expand the node that is **closest to the goal** assuming it will lead to a solution quickly
  - ✓  $f(n) = h(n)$
  - ✓ a.k.a. “Greedy Search”
- Implementation
  - ✓ expand the “most desirable” node into the fringe queue
  - ✓ sort the queue in decreasing order of desirability
- Example: consider the straight-line distance heuristic  $h_{SLD}$ 
  - ✓ Expand the node that appears to be closest to the goal

# CONTINUE...

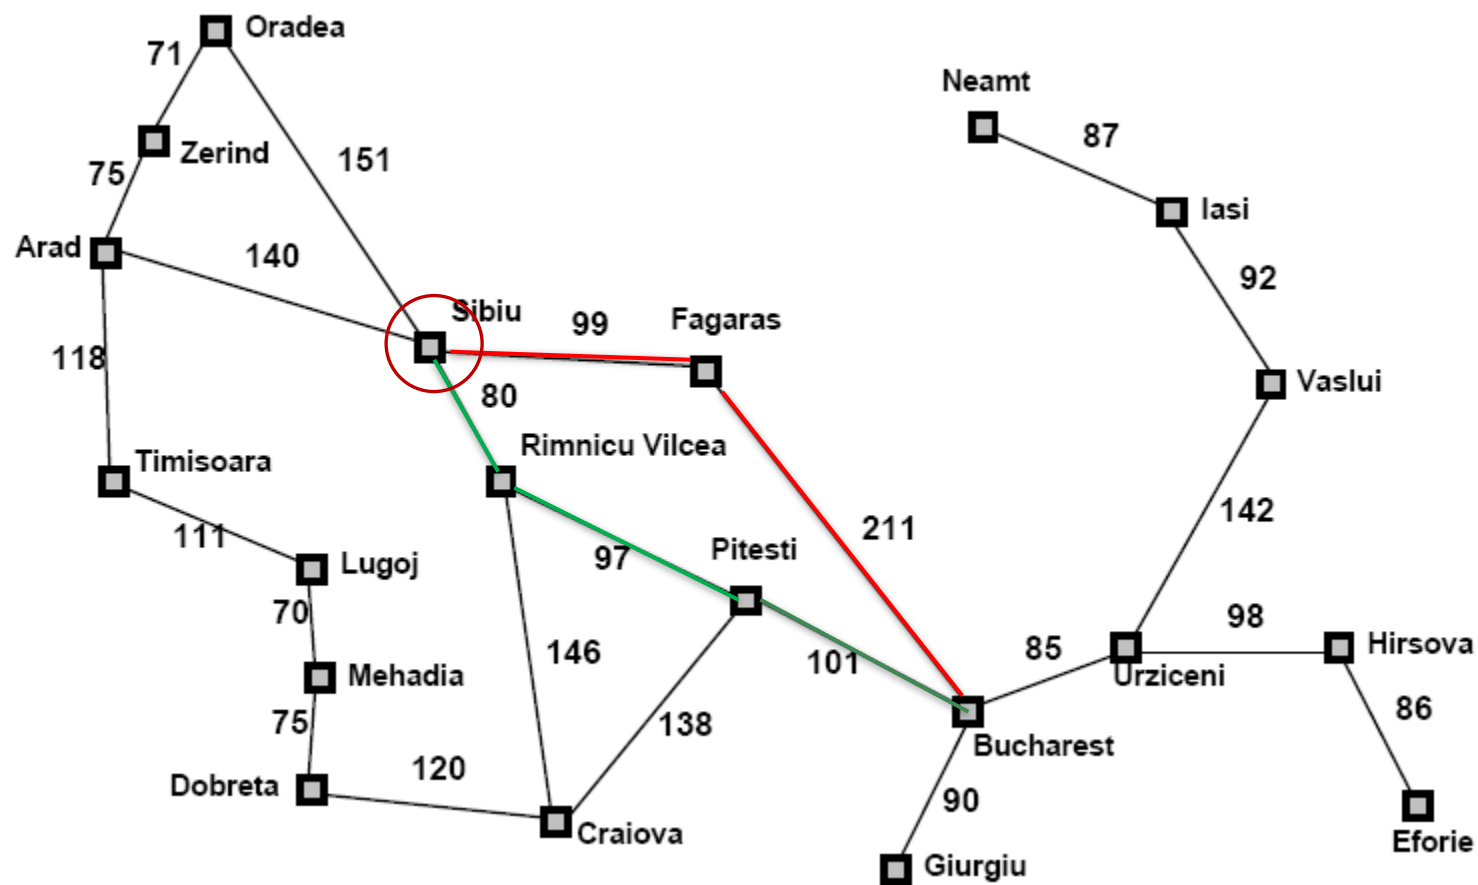


Figure 5: Greedy BFS Example

# CONTINUE...

- Notice that the values of  $h_{\text{SLD}}$  cannot be computed from the problem itself
- It takes some experience to know that  $h_{\text{SLD}}$  is correlated with actual road distances
  - ✓ Therefore a useful heuristic

# CONTINUE...

- **Complete**

- ✓ No, GBFS can get stuck in loops (e.g. bouncing back and forth between cities)

- **Time**

- ✓  $O(bm)$  but a good heuristic can have dramatic improvement

- **Space**

- ✓  $O(bm)$  – keeps all the nodes in memory

- **Optimal**

- ✓ No!

# HEURISTIC FUNCTIONS

- Heuristic Function - function applied to a state in a search space to indicate a likelihood of success if that state is selected:
  - ✓ Heuristic search methods are known as “**weak methods**” because they are dependent on domain specific knowledge.
- The heuristic tells us, how far the state is from the goal state.
- **$h(n)$**

# A\* SEARCH

■ A\* (A star) is the most widely known form of Best-First search

✓ It evaluates nodes by combining  $g(n)$  and  $h(n)$

✓  $f(n) = g(n) + h(n)$

Where,

$g(n)$  = cost so far to reach  $n$

$h(n)$  = estimated cost to goal from  $n$

$f(n)$  = estimated total cost of path through  $n$

# CONTINUE...


- When  $h(n) = \text{actual cost to goal}$ 
  - ✓ Only nodes in the correct path are expanded
  - ✓ Optimal solution is found
- When  $h(n) < \text{actual cost to goal}$ 
  - ✓ Additional nodes are expanded
  - ✓ Optimal solution is found
- When  $h(n) > \text{actual cost to goal}$ 
  - ✓ Optimal solution can be overlooked



# LOCAL SEARCH ALGORITHMS

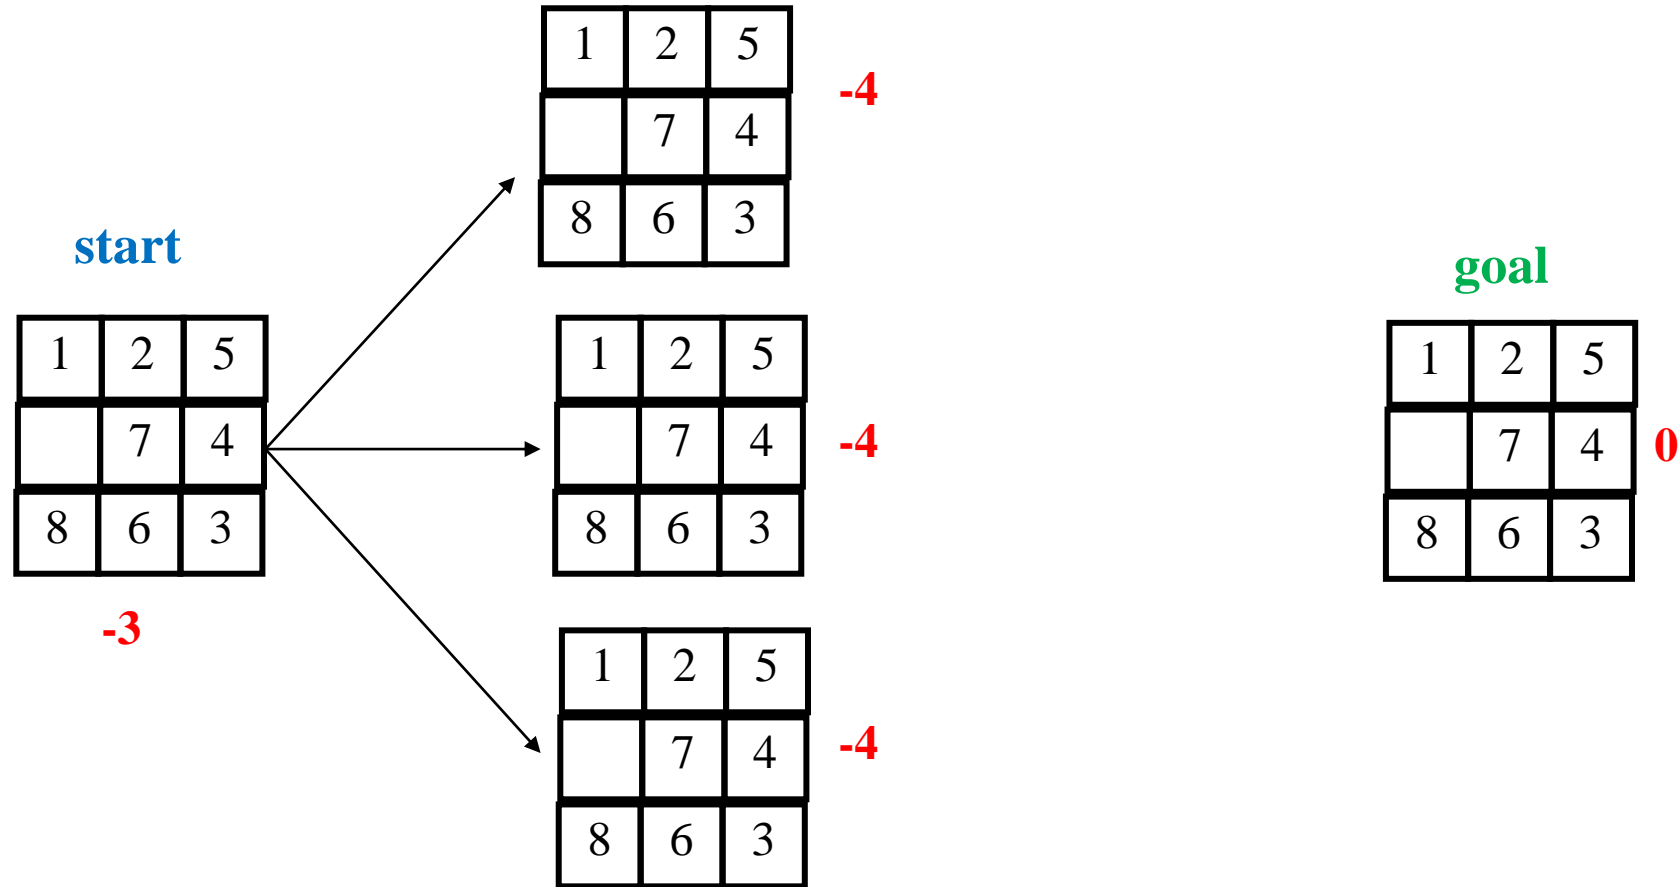
- The concept we have adopted so far is systematic exploration of search space.
- Some problems are path irrelevant type of problems.
- Local search techniques are used here.
- Use single current state and move to neighboring states.
- Advantages:
  - ✓ Use very little memory
  - ✓ Find often reasonable solutions in large or infinite state spaces
- Are also useful for pure optimization problems.
  - ✓ Find best state according to some objective function

# HILL CLIMBING EXAMPLE



$$f(n) = -(\text{number of tiles out of place})$$

# EXAMPLE OF A LOCAL MAXIMUM



# LEARNING HEURISTIC FROM EXPERIENCE

- $h(n)$  is supposed to estimate the cost of a solution beginning from the state at node  $n$ .
- How could an agent construct such a function?
- One of possible solution is learn from the **Experiences**.
- From each optimal solution of a problem,  $h(n)$  can be trained.
- Each example consists of a state from the solution path and the actual cost of the solution from that point.
- From these examples, a learning algorithm can be used to construct a function  $h(n)$  that can predict solution costs for other states that arise during the searching process.

# CONTINUE...

- **Inductive learning** methods work best when supplied with features of a state that are relevant to predicting the state's value, rather than with just the raw state description.
- For example, the feature “**number of misplaced tiles**” might be helpful in predicting the actual distance of a state from the goal.
- Let's call this feature  $x_1(n)$ .
- We could take 100 randomly generated 8-puzzle configurations and gather statistics on their actual solution costs. We might find that when  $X_1(n)$  is 5.
- Given these data, the value of  $X_1$  can be used to predict  $h(n)$ .
- We can even take some other feature, say  $X_2(n)$ . That might be “**number of pairs of adjacent tiles that are not adjacent in goal state**”.
- To combine more than one feature to form a common  $h(n)$ , **linear combination** can be used:  $h(n) = C_1X_1(n) + C_2X_2(n)$ .

# HILL-CLIMBING SEARCH

- The hill-climbing algorithm is simply a loop that continuously moves in the direction of increasing value – that is **Uphill move**.
- It terminates the process when it reaches to peak, where no neighbour has higher value.
- Hill-climbing does not look ahead beyond the immediate neighbours of the current state.
- Hill-climbing can be considered as **Greedy local search**, because it grabs a good neighbour state without thinking about next situations.

# CONTINUE...

**function** HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

*current*  $\leftarrow$  MAKE-NODE(*problem*.INITIAL-STATE)

**loop do**

*neighbor*  $\leftarrow$  a highest-valued successor of *current*

**if** *neighbor*.VALUE  $\leq$  *current*.VALUE **then return** *current*.STATE

*current*  $\leftarrow$  *neighbor*

Figure 6: The Hill-Climbing search algorithm

# CONTINUE...

■ Unfortunately, hill-climbing often gets stuck for the following reasons:

1. Local Maxima
2. Ridges
3. Plateau/Shoulder




Figure 7: Issues with Hill-climbing algorithm



# STIMULATED ANNEALING

- Hill-climbing moves towards goal with **lower value** or **higher value** is guaranteed to be incomplete, as it stuck with local maxima.
- Other side the **random walks** are complete but **extremely inefficient**.
- Stimulated annealing is hybrid formation that achieve efficiency of hill-climbing with completeness of random walks.
- Instead of picking the best move, it picks the random move first. And if the move improves the solution algorithm will accept that move and go further with the hill-climbing approach.
- This was first applied to solve VLSI layout problems in the 1980s.

# LOCAL BEAM SEARCH

- The local beam search keeps track of all required states in the memory.
- It begins with randomly generated states. At each step, all possible successors of all the states are generated.
- If any one is a goal state the algorithm halts, otherwise it selects the best successors from the complete list and repeats.
- In its simple form the algorithm become concentrated on a very small region of the state space, which is very expensive in nature like hill-climbing.
  - The solution is upgrade algorithm – Stochastic Beam Search

# GENETIC ALGORITHM (GA)

- The GA is a variant of beam searching technique, in which successor states are generated by combining two parent states rather than by modifying a single state.
- Similar to beam search, GA begins with a set of randomly generated states, called the **population**.
- Each state is called **individual**, and is represented as a string of digits or alphabets.
- For production of the new states with combination of current state, GA uses **fitness function**. It should return higher values for better state.

# CONTINUE...

**function** GENETIC-ALGORITHM(*population*, FITNESS-FN) **returns** an individual

**inputs:** *population*, a set of individuals

FITNESS-FN, a function that measures the fitness of an individual

**repeat**

*new\_population*  $\leftarrow$  empty set

**for**  $i = 1$  **to** SIZE(*population*) **do**

$x \leftarrow$  RANDOM-SELECTION(*population*, FITNESS-FN)

$y \leftarrow$  RANDOM-SELECTION(*population*, FITNESS-FN)

*child*  $\leftarrow$  REPRODUCE( $x, y$ )

**if** (small random probability) **then** *child*  $\leftarrow$  MUTATE(*child*)

add *child* to *new\_population*

*population*  $\leftarrow$  *new\_population*

**until** some individual is fit enough, or enough time has elapsed

**return** the best individual in *population*, according to FITNESS-FN

---

**function** REPRODUCE( $x, y$ ) **returns** an individual

**inputs:**  $x, y$ , parent individuals

$n \leftarrow$  LENGTH( $x$ );  $c \leftarrow$  random number from 1 to  $n$

**return** APPEND(SUBSTRING( $x, 1, c$ ), SUBSTRING( $y, c + 1, n$ ))

# CONTINUE...

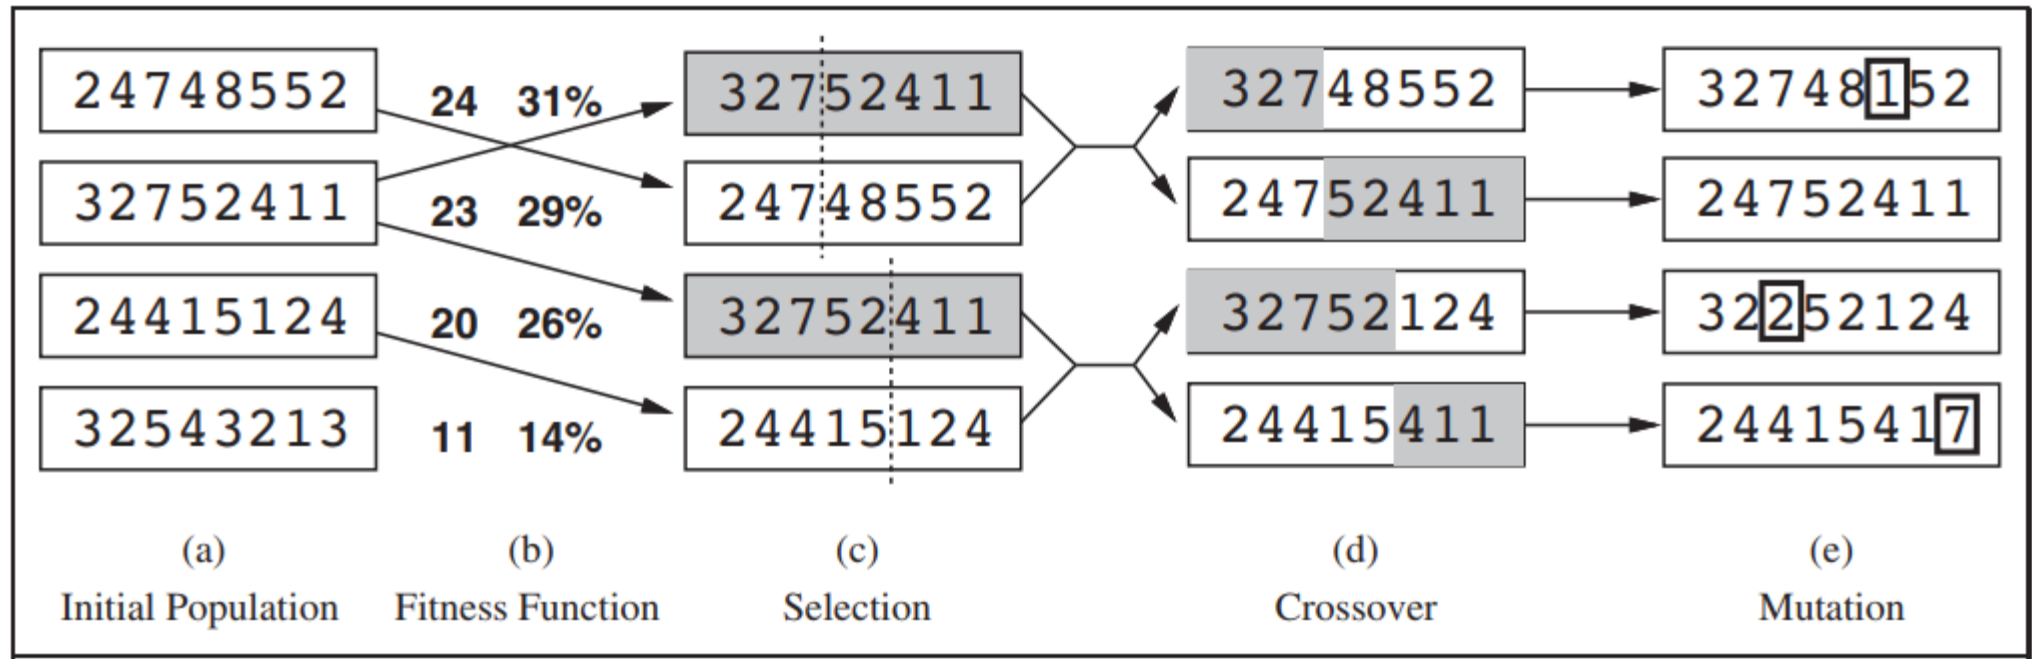


Figure 8: GA explanation with randomly selected 8-digit strings

# CONTINUE...

- Applications of GAs,
  - ✓ Optimization problems
  - ✓ Circuit layout
  - ✓ Job scheduling



# KNOWLEDGE-BASED AGENTS

# KNOWLEDGE-BASED AGENTS

- The human way.
- Agents that can form representations of a complex world using inference to derive new representations about world, and use these new representations to deduce what to do in the task environment, are known as **Logical agents**.



# CONTINUE...

- We learn by the process of reasoning that operates on internal representation of knowledge.
- This approach of human way intelligence is incorporated in AI by **Knowledge-based agents**.
- **Knowledge base** or **KB** is the central component.
- **KB** is a set of sentences. Each sentence is represented in a language called **knowledge representation language**.
- The **TELL** and **ASK** operations for inference – that derives new sentences from old.

# AGENT PROGRAM

- Similar to ordinary agents, they are too having percept as input and they returns an action.
- The agent contains a knowledge base, KB, which may initially have some background knowledge.
- Upon calling agent program does three things.
  1. **TELL** – the KB what it perceives
  2. **ASK** – the KB what actions it should perform
  3. **TELL** – the KB which actions was chosen, and then the agent executes the action

**function** KB-AGENT(*percept*) **returns** an *action*  
**persistent:** *KB*, a knowledge base  
*t*, a counter, initially 0, indicating time

TELL(*KB*, MAKE-PERCEPT-SENTENCE(*percept*, *t*))  
*action* ← ASK(*KB*, MAKE-ACTION-QUERY(*t*))  
TELL(*KB*, MAKE-ACTION-SENTENCE(*action*, *t*))  
*t* ← *t* + 1  
**return** *action*

# CONTINUE...

- **Knowledge level** – we specify only what the agent knows and what its goal are, to fix agent's overall behavior.
- Any knowledge-based agent can be built by simply TELLing it what the agent need to know, in order to achieve the solution in the given task environment.
- It starts with the empty knowledge base, we can TELL sentences one by one until agent knows how to operate in given environment. This is call **declarative approach** to system building.
- The contrast one is **procedural approach**, that encodes the desired behavior directly into agent's program.
- **Declarative vs procedural???**

# THE WUMPUS WORLD PROBLEM

- PEAS description.
- Performance measure:
  - ✓ **+1000 points** for picking up the gold — this is the goal of the agent.
  - ✓ **-1000 points** for dying i.e. entering a square containing a pit or a live Wumpus monster.
  - ✓ **-1** point for each action taken.
  - ✓ **-10** points for using the arrow trying to kill the Wumpus to avoid performing unnecessary actions.

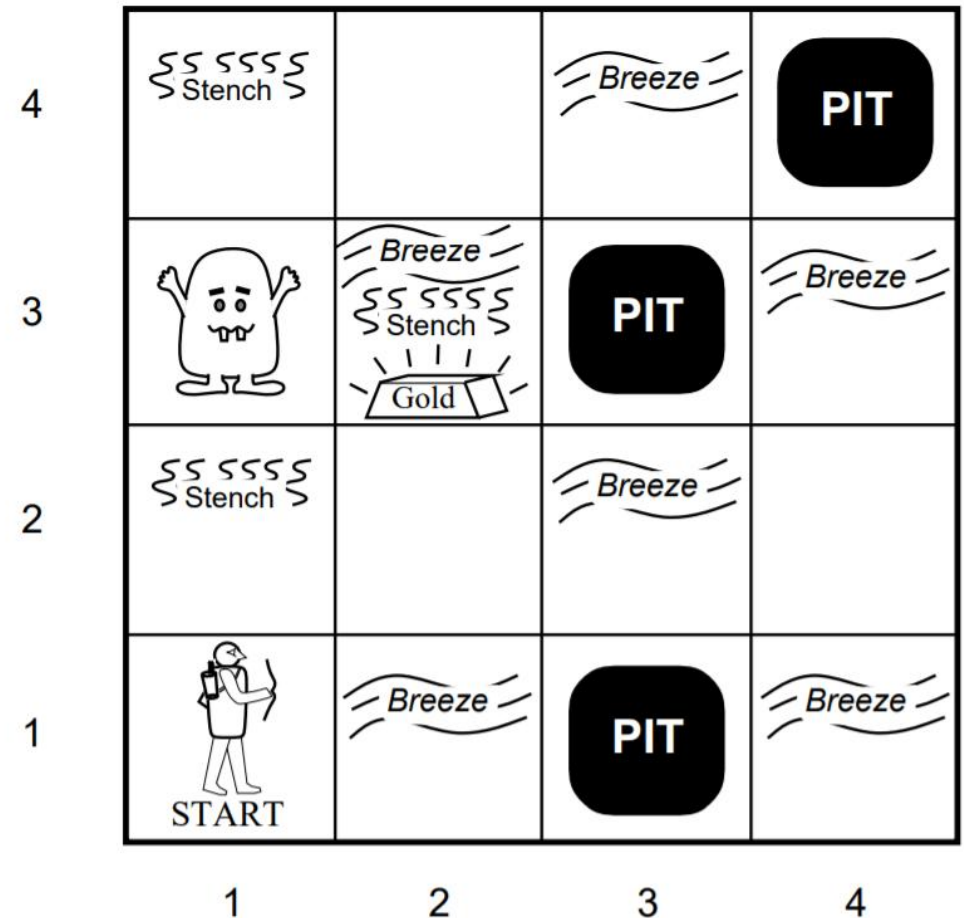



Figure 11: The Wumpus world

# CONTINUE...

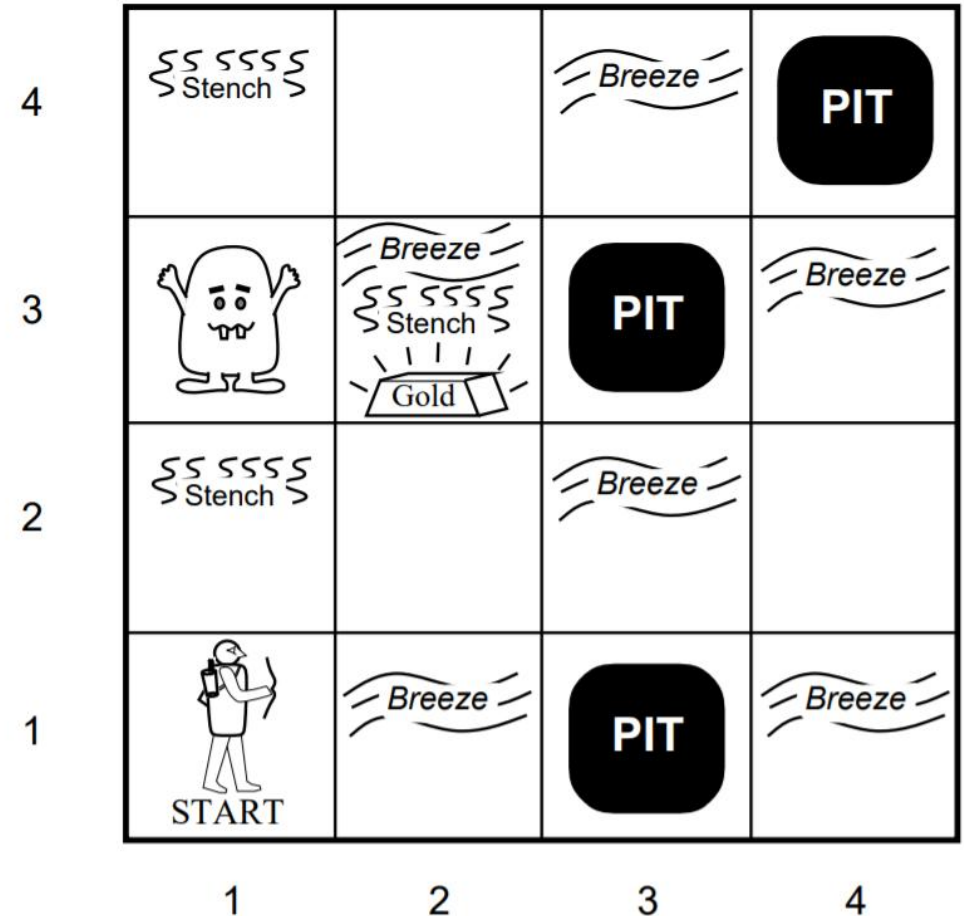
- Environment has A  $4 \times 4$  grid of squares with,
  - ✓ The agent starting from square  $[1, 1]$  facing right.
  - ✓ The gold in one square.
  - ✓ The initially live Wumpus in one square, from which it never moves.
  - ✓ May be pits in some squares.
  - ✓ The starting square  $[1, 1]$  has no Wumpus, no pit, and no gold — so the agent neither dies nor succeeds straight away.



# CONTINUE...

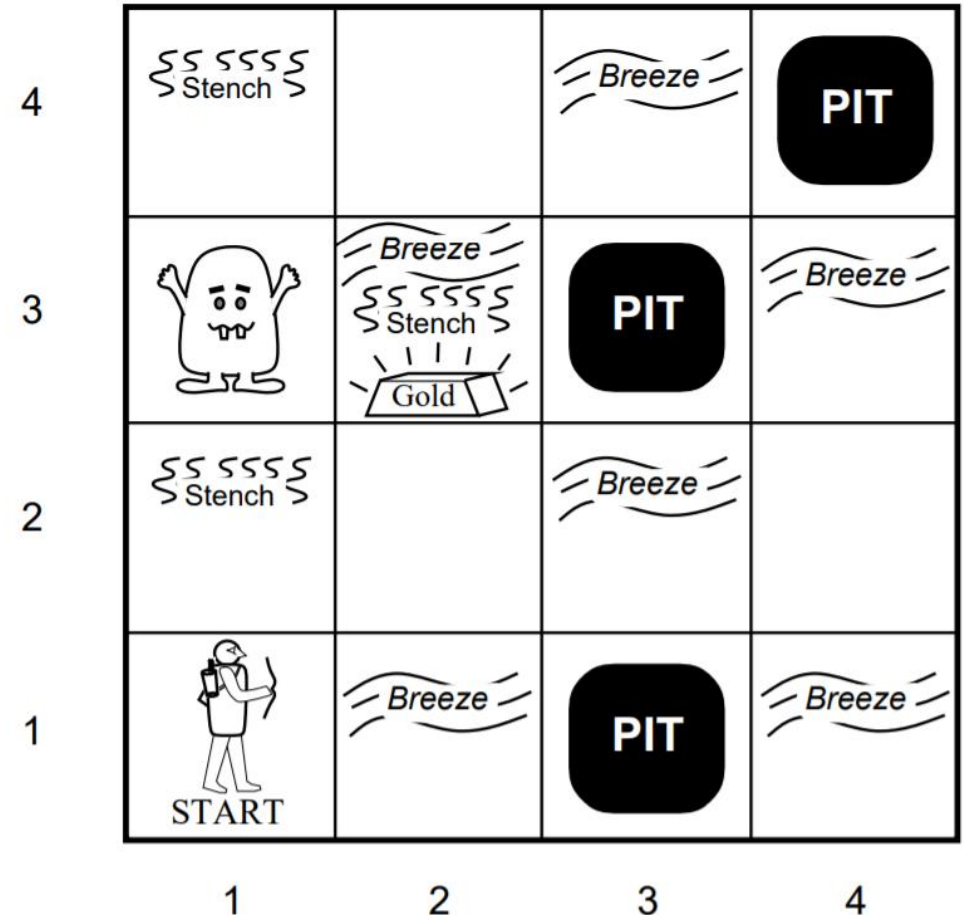
## ■ Actuators:

- ✓ **Turn** 90° left or right.
- ✓ **Walk** one square forward in the current direction.
- ✓ **Grab** an object in this square.
- ✓ **Shoot** the single arrow in the current direction, which flies in a straight line until it hits a wall or the Wumpus.



# CONTINUE...

- **Sensors:** Agent has 5 true/false sensors which reports,
  - ✓ **Stench** when the Wumpus is in an adjacent square — directly, not diagonally.
  - ✓ **Breeze** when an adjacent square has a pit.
  - ✓ **Glitter** when the agent perceives the glitter of the gold in the current square.
  - ✓ **Bump** when the agent walks into an enclosing wall (and then the action had no effect).
  - ✓ **Scream** when the arrow hits the Wumpus, killing it.



- **Top left:** Agent A starts at [1, 1] facing right.
- [1,1] is **OK** as the state is not deadly.
- Agent A gets the percept “**Stench** = **Breeze** = **Glitter** = **Bump** = **Scream** = false”.
- Agent A infers from this percept and  $\beta$  that its both neighbouring squares [1, 2] and [2, 1] are also OK.

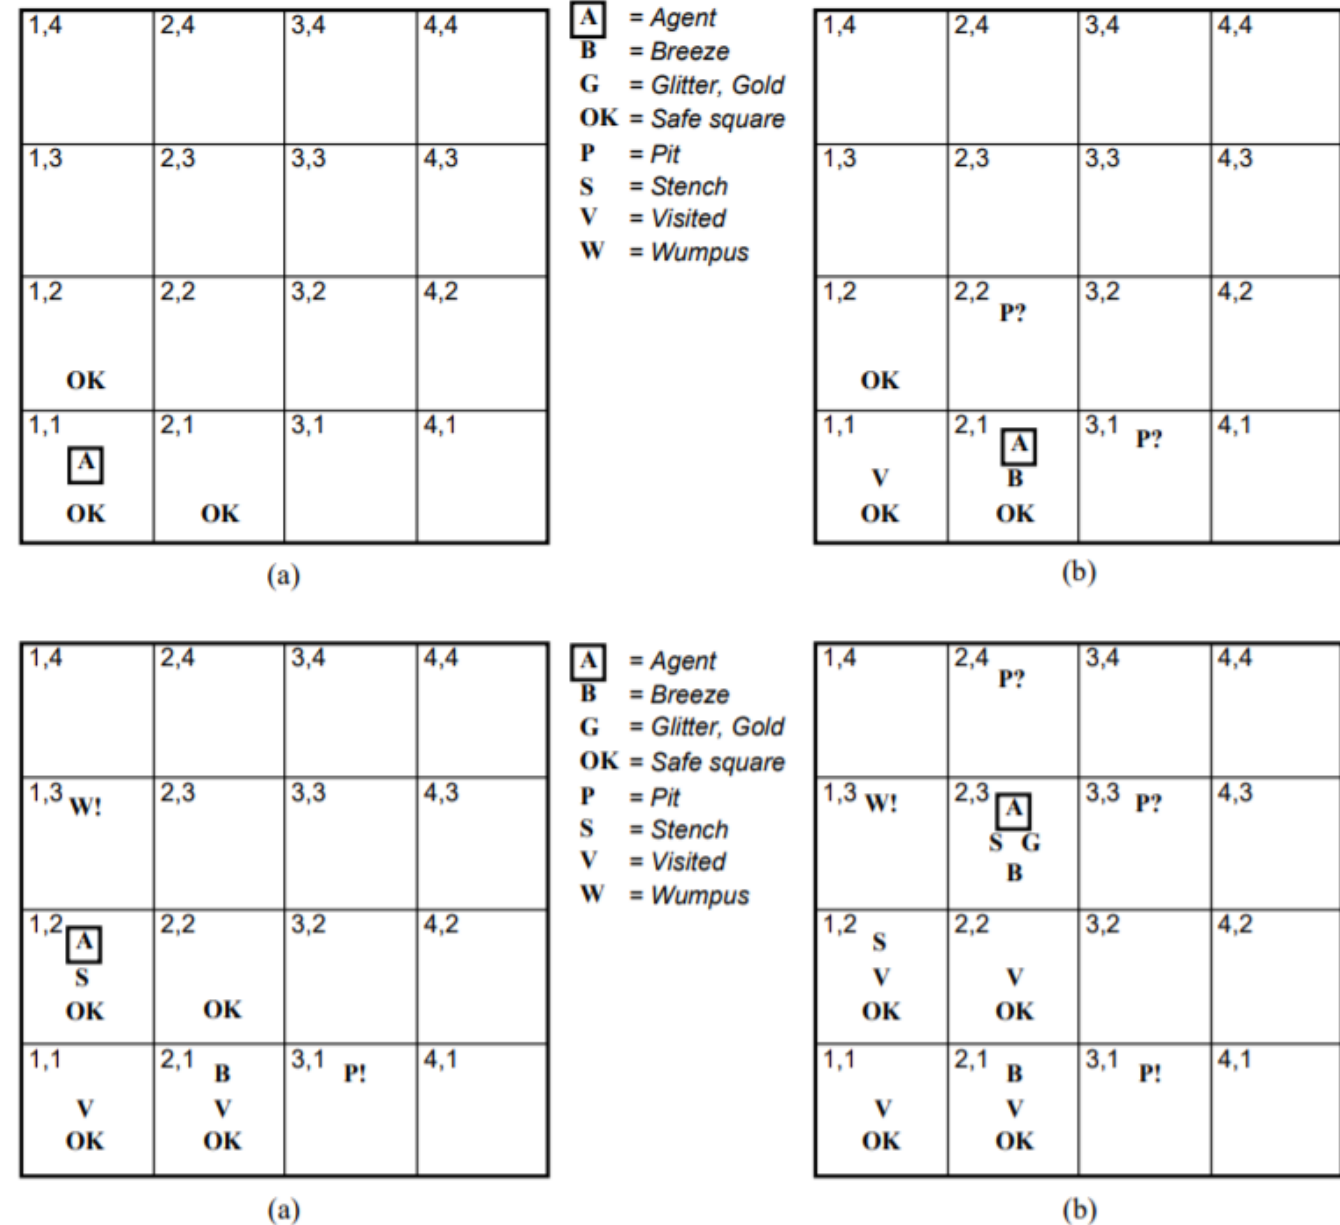


Figure 12: The Wumpus adventure



- **Top Right:** Agent A will only move to **OK** squares.
- Agent A walks into [2, 1], because it is **OK**, and in the direction where agent A is facing, so it is cheaper than the other choice [1, 2]. Agent A also marks [1, 1] Visited.
- Agent A perceives a **Breeze** but nothing else.
- Agent A infers: “At least one of the adjacent squares [1, 1], [2, 2] and [3, 1] must contain a Pit. There is no Pit in [1, 1] by background knowledge  $\beta$ . Hence [2, 2] or [3, 1] or both must contain a Pit.”
- Hence agent A cannot be certain of either [2, 2] or [3, 1], so [2, 1] is a dead end for a Agent A.

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
OK			
1,1	2,1	3,1	4,1
<b>A</b> OK			

(a)

**A** = Agent  
**B** = Breeze  
**G** = Glitter, Gold  
**OK** = Safe square  
**P** = Pit  
**S** = Stench  
**V** = Visited  
**W** = Wumpus

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
OK	P?		
1,1	2,1	3,1	4,1
V OK	<b>A</b> B OK	P?	

(b)

1,4	2,4	3,4	4,4
1,3	W!	2,3	3,3
1,2	<b>A</b> S OK	2,2	3,2
	OK		
1,1	2,1	3,1	4,1
V OK	B V OK	P!	

(a)

**A** = Agent  
**B** = Breeze  
**G** = Glitter, Gold  
**OK** = Safe square  
**P** = Pit  
**S** = Stench  
**V** = Visited  
**W** = Wumpus

1,4	2,4	3,4	4,4
	P?		
1,3	W!	2,3	3,3
	<b>A</b> S G B	P?	
1,2	S V OK	2,2	3,2
	V OK		
1,1	2,1	3,1	4,1
V OK	B V OK	P!	

(b)

- **Bottom Left:** Agent A has turned back from the dead end [2, 1] and walked to examine the other OK choice [1, 2] instead.
- Agent A perceives a Stench but nothing else.
- Agent infers that the Wumpus is in an adjacent square. It was not in [1,1], not in [2,1] either, else it would have sensed the stench in [2,1]. Hence, it is in [1,3].
- There is not breeze here in [1,2] so there is no pit in any adjacent square, i.e. [2,2].
- It finally infers [2,2] as OK.

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
OK			
1,1	2,1	3,1	4,1
<b>A</b> OK	OK		

(a)

**A** = Agent  
**B** = Breeze  
**G** = Glitter, Gold  
**OK** = Safe square  
**P** = Pit  
**S** = Stench  
**V** = Visited  
**W** = Wumpus

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
OK	P?		
1,1	2,1	3,1	4,1
V OK	<b>A</b> B OK	P?	

(b)

1,4	2,4	3,4	4,4
1,3	W!	2,3	3,3
1,2	<b>A</b> S OK	2,2	3,2
	OK		
1,1	2,1	3,1	4,1
V OK	B V OK	P!	

(a)

**A** = Agent  
**B** = Breeze  
**G** = Glitter, Gold  
**OK** = Safe square  
**P** = Pit  
**S** = Stench  
**V** = Visited  
**W** = Wumpus

1,4	2,4	3,4	4,4
	P?		
1,3	W!	2,3	3,3
	<b>A</b> S G B	P?	
1,2	S V OK	2,2	3,2
	V OK		
1,1	2,1	3,1	4,1
V OK	B V OK	P!	

(b)

- **Bottom Right:** Agent A walks to the only unvisited OK choice [2, 2].
- There is no Breeze here, and since the square of the Wumpus is now known too, [2, 3] and [3, 2] are OK too.
- **Agent A walks into [2, 3] and senses the Glitter there, so he grabs the gold and succeeds.**

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
OK			
1,1	2,1	3,1	4,1
A OK	OK		

(a)

A = Agent  
 B = Breeze  
 G = Glitter, Gold  
 OK = Safe square  
 P = Pit  
 S = Stench  
 V = Visited  
 W = Wumpus

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2 P?	3,2	4,2
OK			
1,1	2,1	3,1 P?	4,1
V OK	A B OK		

(b)

1,4	2,4	3,4	4,4
1,3 W!	2,3	3,3	4,3
1,2	2,2	3,2	4,2
A S OK	OK		
1,1	2,1	3,1 P!	4,1
V OK	B V OK		

(a)

A = Agent  
 B = Breeze  
 G = Glitter, Gold  
 OK = Safe square  
 P = Pit  
 S = Stench  
 V = Visited  
 W = Wumpus

1,4	2,4 P?	3,4	4,4
1,3 W!	2,3	3,3 P?	4,3
	A S G B		
1,2	2,2	3,2	4,2
S V OK	V OK		
1,1	2,1	3,1 P!	4,1
V OK	B V OK		

(b)



# LOGICS

# LOGIC AND INFERENCE

- For any agent to act in a meaningful manner, it must have some kind of symbolic representations.
- By intelligent activity we means not just optimization or survival, but something that involves awareness of goals, awareness of situation and informed decision making.
- **Formal logic:**
  - It is primary machinery for realizing the reasoning.
  - **Given a set of completely true statements**, the machinery determines what **other sentences** are argued to be true.

# CONTINUE...

- Formal logic example:

FROM: "All kings are brave."

AND: "Krishna is a king."

INFER: "Krishna is brave."

FROM: "All detectives are rich."

AND: "Byomkesh is detective."

INFER: "Byomkesh is rich."

} Sounds falsy inference

# ENTAILMENT

- In some way all humans have instinct curiosity, a desire to know the truth about something or the other.
- To a large extent we rely upon our senses.
- One another source that we unknowingly utilize is **LOGIC**.  
We ignore the rumours and other speculations that sounds unreliable.
- In logic we are focused to deal with only true statements.
- So, given a collection of true statements, also called **premises**, we are interested in knowing “what other/new statements are logically entailed/made true with the help of premises.”

# PROOFS

- We know, entailment is concerned with true statements.
- To determine, whether the statement is true or not true is not straightforward.
- Logic uses concept of proofs. Which is made up of a sequence of inference steps.
- Each inference step allow us to add one more new sentence to the existing set of sentences.
- Each inference step is based on *rule of inference*.
- Example:
  - FROM: "ALL x's are y's"
  - AND: "k is x"
  - INFER: "k is y"



# SOUNDNESS AND COMPLETENESS

- A logic is said to be sound if it produces only true statements or in other words if it does not produce false statements.
- A logic is complete if it produces all true statements.
- Formally,
  - If  $P$ , provable statement is a subset of  $T$ , true statement then logic is SOUND.
  - If  $T$  is subset of  $p$ , then logic is COMPLETE.

# PRINCIPLES BEHIND LOGICS

- The actual syntax in which the logical statements are written as formulas is just a “print-out” of the underlying data structure.
- This underlying data structure is the parse tree of the formula.
- From a programming viewpoint, implementations of logical inference create from existing such structures new ones in ways which respect their **meaning**.
- This **meaning** or **semantics** of a statement defines whether it is true or false in the given possible world.
- In our agent viewpoint, possible world = possible state of the environment.

# CONTINUE...

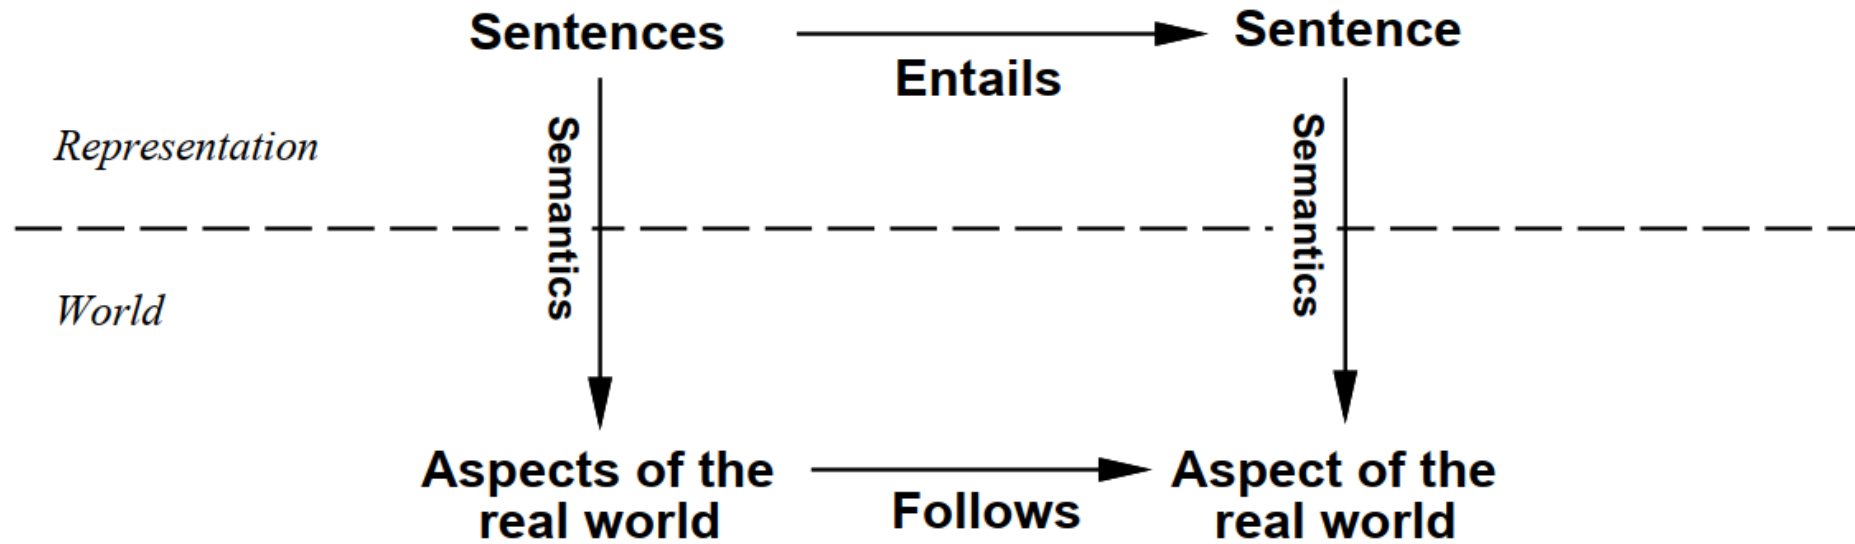


Figure 13: The world of representation

# PROPOSITIONAL LOGIC

- The smallest “addressable unit” of propositional logic is a whole sentence which can be either true or false.
- That is, any expression  $\eta$  such that
  - ✓ “Is it the case that  $\eta$ ?”
- Asking this about  $\eta$  = “Manoj is devotee of Krishna”
- On the other hand, e.g. neither “Is it the case that devotee?” nor “Is it the case that Krishna?” is a meaningful question.
- Note that “Is it the case that Manoj devotee?” is another meaningful question — but this is a different sentence.

# CONTINUE...

- The **syntax** of propositional logic defines the allowable sentences.
- The **atomic sentences** consist of a single **proposition symbol**.
- Each such symbol stands for a proposition that can be true or false. We use symbols that start with an uppercase letter and may contain other letters or subscripts, for example: P, Q, R, W<sub>1,3</sub> and North.
- The names are arbitrary but are often chosen to have some mnemonic value—we use W<sub>1,3</sub> to stand for the proposition that the Wumpus is in [1,3].

# CONTINUE...

- **Complex sentences** are constructed from simpler sentences, using parentheses and **logical connectives**.
- There are five common connectives to be used:
  1.  $\neg$  (**not**). A sentence such as  $\neg W_{1,3}$  is called the **negation** of  $W_{1,3}$ .
  2.  $\wedge$  (**and**). A sentence whose main connective is  $\wedge$ , such as  $W_{1,3} \wedge P_{3,1}$ , is called a **conjunction**.
  3.  $\vee$  (**or**). A sentence using  $\vee$ , such as  $(W_{1,3} \wedge P_{3,1}) \vee W_{2,2}$ , is a **disjunction** of the **disjuncts**  $(W_{1,3} \wedge P_{3,1})$  and  $W_{2,2}$ .
  4.  $\Rightarrow$  (**implies**). A sentence such as  $(W_{1,3} \wedge P_{3,1}) \Rightarrow \neg W_{2,2}$  is called an **implication**.
  5.  $\Leftrightarrow$  (**if and only if**). The sentence  $W_{1,3} \Leftrightarrow \neg W_{2,2}$  is a **biconditional**.

# CONTINUE...

$$\begin{aligned} \textit{Sentence} &\rightarrow \textit{AtomicSentence} \mid \textit{ComplexSentence} \\ \textit{AtomicSentence} &\rightarrow \textit{True} \mid \textit{False} \mid P \mid Q \mid R \mid \dots \\ \textit{ComplexSentence} &\rightarrow (\textit{Sentence}) \mid [\textit{Sentence}] \\ &\mid \neg \textit{Sentence} \\ &\mid \textit{Sentence} \wedge \textit{Sentence} \\ &\mid \textit{Sentence} \vee \textit{Sentence} \\ &\mid \textit{Sentence} \Rightarrow \textit{Sentence} \\ &\mid \textit{Sentence} \Leftrightarrow \textit{Sentence} \end{aligned}$$

OPERATOR PRECEDENCE :  $\neg, \wedge, \vee, \Rightarrow, \Leftrightarrow$

Figure 14: Formal grammar of proposition logic with BNF (Backus-Naur Form) notations

# BNF GRAMMAR

- The BNF grammar by itself is **ambiguous**; a sentence with several operators can be parsed by the grammar in multiple ways.
- **To eliminate the ambiguity** we define a **precedence** for each operator.
- The **“not” operator** ( $\neg$ ) has the **highest precedence**, which means that in the sentence  $\neg A \wedge B$  the  $\neg$  binds most tightly, giving us the equivalent of  $(\neg A) \wedge B$  rather than  $\neg(A \wedge B)$ .



# PROPOSITIONAL THEOREM PROVING

- Applying rules of inference directly to the sentences in our knowledge base to construct a proof of the desired sentence without consulting models, called theorem proving.
- If the number of models is large but the length of the proof is short, then theorem proving can be more efficient.
- **Logical equivalence**: two sentences  $\alpha$  and  $\beta$  are logically equivalent if they are true in the same set of models. i.e.  $\alpha \equiv \beta$ .
- Any two sentences  $\alpha$  and  $\beta$  are equivalent only if each of them entails the other:
  - ✓  $\alpha \equiv \beta$  if and only if  $\alpha \models \beta$  and  $\beta \models \alpha$
- **Validity**: A sentence is valid if it is true in *all* models. For example, the sentence  $P \vee \neg P$  is valid. Valid sentences are also known as **tautologies**—they are *necessarily* true.

# CONTINUE...

- **Satisfiability**: A sentence is satisfiable if it is true in, or satisfied by, *some* model.

# REFERENCES

- [1] Artificial Intelligence – A Modern Approach, 3<sup>rd</sup> Edition, Stuart Russell and Peter Norvig, Pearson Publication
- [2] <https://www.humanbrainfacts.org/human-brain-memory.php>
- [3] <https://operativeneurosurgery.com/doku.php?id=hippocampus>