

UNIT - 3

STATISTICAL LEARNING

By Vishvajit Bakrola

MACHINE LEARNING

- Is a study and implementation of algorithms, in order to program computer to optimize a performance criterion using example data or past experience.
- Learning will become essential when,
 - ✓ Human expertise does not exist
 - ✓ Humans are unable to explain their expertise
 - ✓ Solution changes in time
 - ✓ Solution needs to be adapted to particular cases

MACHINE LEARNING - DEFINITION

- “Machine learning is a study of computer algorithms that allow computer programs to automatically improve through experience.”
- “The field of study that gives computers the ability to learn without being explicitly programmed.”
- “A computer program is said to learn from experience E with respect to some class of tasks T and the performance measure P , if its performance at task T , as measured by P , improves with experience E .”

DEFINING THE LEARNING TASK

□ Improve on task T , with respect to performance metric P , based on experience E

T : Playing board game

P : Percentage of games won against an arbitrary opponent


E : Playing practice games against itself

T : Recognizing hand-written characters

P : Percentage of characters correctly classified

E : Database of human-labeled images of handwritten characters

RELATED FIELDS



LEARNING ASSOCIATION

Basket Analysis – Finding association between products bought by customers.

- *If there is a customer who buy X typically also buy Y, and if there is a customer who buys X and does not buy Y, he or she is a potential Y customer.*
- This can be helpful for cross-selling.

Association Rule – Learning a conditional probability of the form $P(Y | X)$, *Y is the product we would like to conditioned on X, which is a product which we know that the customer has already purchased.*

$P(\text{Biscuits} | \text{Tea}) = 0.85$ → We can define the rule that “85 percent of customers who buy **Tea** also buys **Biscuits**”

This can possibly $P(Y | X, D)$, D:Customer Attributes

TYPES OF LEARNING

1. Supervised Learning

- Training data + Desired outputs(labels) are given

2. Unsupervised Learning

- Training data given without desired outputs

3. Semi-supervised Learning

- Training data + Partial desired outputs available


4. Reinforcement Learning

- Only collects rewards from sequence of actions

CLASSIFICATION

- When we need to assign individual unique label to classes/groups.
- Example: Credit scoring
- Differentiating between **low-risk** and **high-risk** customers from their *income* and *savings*.

Discriminant: IF *income* $> \theta_1$ AND *savings* $> \theta_2$ THEN **low-risk** ELSE **high-risk**




APPLICATIONS OF CLASSIFICATION

- ❑ **Face recognition:** Pose, lighting, occlusion (glasses, beard), make-up, hair style
- ❑ **Hand-written character recognition:** Different handwriting styles
- ❑ **Medical diagnosis:** From symptoms to illnesses
- ❑ **Others:** E-mail spam, Image segmentation, Speech recognition, Genetic sequencing classification, etc.

REGRESSION

- Belongs to prediction types of problems, and their output is a number.
- Let, X is some attributes of a car and Y be the price of a car.
- Surveying the previous such transaction we can collect the training data. Fitting appropriate machine learning model gives us predicted price of car given the attributes.
- We use linear model (*simple equation of linear line i.e. $Y=mx+c$*)
- We get, **$Y = WX + W_0$**

CONTINUE...



CONTINUE...

- The model is linear as W and W_0 are the parameters optimized for best fit to the training data.
- If our data is type of non-linear, we can make our model quadratic or higher order polynomial.

$$Y = W_2X^2 + W_1X + W_0$$

SIMPLE LINEAR REGRESSION

- We can predict score of one variable from the score of second variable.
- The variable(s) that we are predicting is **criterion variable** and we will refer it as Y.
- The variable we are basing upon for making predictions in called **predictor variable**.
- The general equation we get is,

$$Y = \beta_0 + \beta_1 X$$

Dependent variable

Intercept
Value of Y
when $x=0$

Slop

Explanatory variable

CONTINUE...

- Linear regression is an approach for modelling the relationship between a scalar dependent variable and one or more explanatory variables.
- The specific case of one explanatory variable is called **Simple Linear Regression**.
- NOTE:
 - ✓ When we use X to predict Y , called **Relationship estimation**
 - ✓ To estimate effect of X on Y , called **Forecast**

EXAMPLE

Student ID	x_i	y_i	$(x_i - \bar{X})$	$(y_i - \bar{Y})$	$(x_i - \bar{X})^2$	$(y_i - \bar{Y})^2$	
1	98	85	20	8	400	64	
2	85	95	7	18	49	324	
3	80	70	2	-7	4	49	
4	70	65	-8	-12	64	144	
5	60	70	-18	-7	324	49	
Sum	390	385					
Mean	78	77					

CONTINUE...

$$b_1 = \Sigma [(x_i - \bar{x})(y_i - \bar{y})] / \Sigma [(x_i - \bar{x})^2]$$

$$\text{So, } b_1 = 494/730$$

Finally **$b_1 = 0.644$**

CONTINUE...

- Using regression coefficient b_1 , we can solve for regression slope b_0 .

$$\begin{aligned} b_0 &= y - b_1 * \mathbf{X} \\ &= 77 - (0.644) * 78 \\ &= \mathbf{26.768} \end{aligned}$$

CONTINUE...

Now we have, **b1 = 0.644** and **b0 = 26.768**


We get,

$$Y = 26.768 + 0.644X$$

$$Y = 26.768 + 0.644 * 80$$

$$Y = 78.288$$

3	80	70
---	----	----



This is called **extrapolation**.

CONTINUE...

- Whenever, we use the regression equation we need to understand how well the equation fits the data.
- One solution to achieve these is to find **coefficient of determination**.

$$R^2 = \left\{ \left(\frac{1}{N} \right) * \sum \frac{[(x_i - X) * (y_i - Y)]}{(\sigma_x * \sigma_y)} \right\}^2$$

N is Number of observations used to fit the model

σ_x is Standard deviation of x

σ_y is Standard deviation of y

CONTINUE...

$$R^2 = \left\{ \left(\frac{1}{N} \right) * \sum \frac{[(x_i - X) * (y_i - Y)]}{(\sigma_x * \sigma_y)} \right\}^2$$

$$R^2 = \left\{ \left(\frac{1}{5} \right) * \frac{470}{(12.083 * 11.255)} \right\}^2$$

$$R^2 = \left(\frac{94}{135.632} \right)^2$$

$$R^2 = (0.693)^2$$

$$R^2 = \mathbf{0.48}$$

Indicates 48% of variation in statistics grades (DV) can be explained by IV

CLASSIFICATION

- The understanding of boundary conditions that can be use to determine each target class in training data.
- Once the boundary conditions are determined the next task is to predict the target class.
- The process formally known as **classification**.
- Examples:
 - ✓ Analysis of student data, whether he/she will buy a laptop or not. (**Target class: Yes or No**)
 - ✓ Classifying fruits using associated features like color, shape, size, weight, etc. (**Target class: Name of the Fruit**)
 - ✓ Student classification using features of student-uniform. (**Target class: Name of institute**)

CONTINUE...

- **Classifier:** A technique or an algorithm that maps input data to specific category.
- **Classification model:** Retrieve some meaningful conclusion from the input data given during training in order to predict the class labels.
- **Feature:** Is an individual and unique measurable property being observed from the data.
- **Binary classification:** The task of classification having possible outcomes.
- **Multi-class classification:** The task of classification having more than two outcomes.
- **Multi-label classification:** The task of classification having more than two class labels. One or more class labels may be predicted for each example.

GENERAL TYPES OF CLASSIFICATION ALGORITHMS

☐ Linear classification

- Logistic Regression
- Naïve Bayes Classifier
- Linear discriminant

☐ SVMs

☐ Quadratic classifiers

☐ Decision Trees

☐ Neural Networks

☐ LVQs


SUPPORT VECTOR MACHINES

- Linear classifier is there but which hyperplane to select?
- Lots of possible solutions for a , b , c .
- Some methods find a separating hyperplane, but not the optimal one.

Support Vector Machine (SVM) finds an optimal solution.


- Maximizes the distance between the hyperplane and the “difficult points” close to decision boundary
- One intuition: if there are no points near the decision surface, then there are no uncertain classification decisions.

This line represents the decision boundary:
 $ax + by - c = 0$




CONTINUE...

- If we have to place a fat separator between classes, we have less choices, and so the capacity of the model has been decreased.



CONTINUE...

- SVMs **maximize the margin** around the separating hyperplane.
- The decision function is fully specified by a subset of training samples, the **support vectors**.
- Solving SVMs is a **quadratic programming problem**.
- Seen by many as the most successful text classification method.



CONTINUE...


- **w**: decision hyperplane normal vector
- **x_i**: data point *i*
- **y_i**: class of data point *i* (+1 or -1)
- Classifier is:
$$\mathbf{f}(\mathbf{x}_i) = \mathbf{w}^T \mathbf{x}_i + b$$
- Functional margin of **x_i** is: $y_i (\mathbf{w}^T \mathbf{x}_i + b)$
 - But note that we can increase this margin simply by scaling **w**, **b**....
- Functional margin of dataset is twice the minimum functional margin for any point
 - The factor of 2 comes from measuring the whole width of the margin

CONTINUE...

- Distance from example to the separator is

$$r = y \frac{\mathbf{w}^T \mathbf{x} + b}{\|\mathbf{w}\|}$$

- Examples closest to the hyperplane are **support vectors**.
- Margin** ρ of the separator is the width of separation between support vectors of classes.
- Dotted line $\mathbf{x}' - \mathbf{x}$ is parallel to **decision boundary**.



CONTINUE...

- Following two constraints follow for a training set $\{(\mathbf{x}_i, y_i)\}$.

$$\mathbf{w}^T \mathbf{x}_i + b \geq 1 \quad \text{if } y_i = 1$$

$$\mathbf{w}^T \mathbf{x}_i + b \leq -1 \quad \text{if } y_i = -1$$

- Then, since each example's distance from the hyperplane is $r = y \frac{\mathbf{w}^T \mathbf{x} + b}{\|\mathbf{w}\|}$
- The margin is: $r = \frac{2}{\|\mathbf{w}\|}$


DISCRIMINANT FUNCTIONS

- A discriminant is a function that takes an input vector x and assigns it to one of k -classes, denoted by C_k .
- $K > 2$ is for multiclass problem.
- The simplest representation of a linear discriminant function is obtained by taking a linear function of input vector so that,

$$y(x) = \underset{\substack{\uparrow \\ \text{Weight Vector}}}{W^T} \cdot X + \underset{\substack{\uparrow \\ \text{Bias}}}{W_0}$$


CONTINUE...

- An input vector x is assigned to class, $C_1 = y(x) \geq 0$ and $C_2 = \textit{otherwise}$
- The corresponding decision boundary is therefore defined as $y(x) = 0$.
- $y(x) = 0$ corresponds to a $(D-1)$ dimensional hyperplane within the D -dimensional input space.




CONTINUE...

- It is extension of linear discriminant to $k > 2$ classes.
- Use of $k-1$ classifiers where, each solves a two class problem of separating points in a particular class C_k from points not in that class, known as a **“one-versus-the-rest” classifier**.



CONTINUE...

- The examples involving three classes, where this approach leads to regions of input space that are ambiguously classified.
- An alternative is to introduce $k(k-1)/2$ binary discriminant function. One for every possible pair of classes. This is known, **“one-versus-the-one” classifier**.



VARIANCE AND COVARIANCE

- Variance and Covariance are a measure of the “spread” of a set of points around their center of mass (mean).
- Variance – measure of the deviation from the mean for points in one dimension.
- Covariance as a measure of how much each of the dimensions vary from the mean with respect to each other.
- Covariance is measured between 2 dimensions to see if there is a relationship between the 2 dimensions e.g. number of hours studied & marks obtained.
- **The covariance between one dimension and itself is the variance**
- Used to find relationship between dimensions in the high dimensional data sets, where visualization is difficult.

CONTINUE...

- Covariance between x and y,





$$(x, y) = \frac{\sum_{i=1}^n (\bar{X}_i - X)(\bar{Y}_i - Y)}{(n - 1)}$$

CONTINUE...

- If we have a 3-dimensional data set (x, y, z) , then we first measure the covariance between the x and y dimensions, then y and z dimensions, and then x and z dimensions.
- A **positive value** of covariance indicates **both dimensions increase or decrease together** e.g. as the number of hours studied increases, the marks in that subject increase.
- If **covariance is zero**: the **two dimensions are independent** of each other e.g. heights of students vs the marks obtained in a subject
- A **negative value** indicates while **one increases the other decreases, or vice-versa**.

DIMENSIONALITY REDUCTION

- Increasing the number of features will not always improve classification accuracy – **Curse of Dimensionality**.
- In practice, the inclusion of more features might actually lead to **worse** performance.
- The number of training examples required increases **exponentially** with dimensionality d (i.e., k^d).



CONTINUE...

- **Visualization** – Projections of higher dimensional data to 2D or 3D.
- **Removal of noise** – Removing noise gives the clarity in data and positive impacts on accuracy.
- **Compression of data** – Ultimately leads to efficient storage and easy data retrieval.

CONTINUE...

- Our motto is to choose an optimum set of features of lower dimensionality to **improve** classification accuracy.
- **Feature extraction:** To find a set of **new** features (i.e., through some mapping **f()**) from the **existing** features.

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ x_N \end{bmatrix} \xrightarrow{f(\mathbf{x})} \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \cdot \\ \cdot \\ y_K \end{bmatrix}$$

The mapping $f()$
could be linear or
non-linear.
Here, $K < N$

CONTINUE...

- **Feature selection:** chooses a subset of the original features.
- **Linear** combinations are particularly attractive because they are simpler to compute and analytically tractable.
- Commonly used linear feature extraction methods:
 - **Principal Components Analysis (PCA):** Seeks a projection that **preserves** as much **information** in the data as possible.
 - **Linear Discriminant Analysis (LDA):** Seeks a projection that **best discriminates** the data.
 - Few other methods are - **Projection Pursuit, Independent Component Analysis or ICA, Isomap, Locally Linear Embedding or LLE, etc.**

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \\ x_N \end{bmatrix} \rightarrow \mathbf{y} = \begin{bmatrix} x_{i_1} \\ x_{i_2} \\ \cdot \\ \cdot \\ x_{i_K} \end{bmatrix}$$

LINEAR DISCRIMINANT ANALYSIS

- Is dimensionality reduction technique used as a pre-processing.
- As we know that the main goal of dimensionality reduction is to remove redundant and dependent features by transforming them in lower dimensions.
- LDA is supervised technique, as it takes labels into consideration.
- In LDA, we first calculate the separability between two classes and then the distance between mean and sample of each class.

CONTINUE...

- Step – 1: To calculate distance between mean of different classes.

$$S_b = \sum_{i=1}^{\mathcal{G}} N_i (\bar{x}_i - \bar{x})(\bar{x}_i - \bar{x})^T$$

CONTINUE...

- Step – 2: To calculate distance between mean and sample of each particular class.

$$S_w = \sum_{i=1}^{\mathcal{G}} (N_i - 1) S_i = \sum_{i=1}^{\mathcal{G}} \sum_{j=1}^{N_i} (x_{i,j} - \bar{x}_i)(x_{i,j} - \bar{x}_i)^T$$

CONTINUE...

Step – 3: To project the lower dimensional space which maximizes the variance between classes and minimizes the in-class variance.

$$P_{lda} = \arg \max_P \frac{|P^T S_b P|}{|P^T S_w P|}$$

MAXIMUM LIKELIHOOD ESTIMATION (MLE)

- Statistic can be any value that is calculated from a given sample.
- In the process of inference with statistic, we make decision using the information provided by a sample.
- One of the approaches is parametric approach, where we assume that the sample is drawn from some distribution that follows our known model.
- It can be drawn from small number of parameters, and once those parameters are known the whole model is known to us.

MAXIMUM LIKELIHOOD ESTIMATION (MLE)

- MLE is used for estimation of critical **parameters** of our model.
- In machine learning we are referring a model as a process that ultimately results us the data that are observed.
- I.e. object classification with object properties, spend on advertisement vs. revenue generated, etc.
- All such models is having its own set of parameters, that ultimately defines what exactly the model is.

CONTINUE...

- From linear equation $y=mx+c$, where x is the spending on advertisement and y will be revenue generated.
- m and c will be the parameters for this model. We understand the different values of this model gives different association between values of x and y .
- Ultimately parameters are defining how model will react on the data.
- MLE is a method that determines the values of such parameters of our model.
- Here, MLE returns parameters values that maximize the likelihood of our sample.

CONTINUE...

- **Bernoulli distribution** – There are two outcomes that whether the event will occur or not. i.e. **any given instance is a positive example of a class or not.**
- **Multinomial distribution** – It is generalization of Bernoulli distribution, where instead of two states the outcome of a random event is one of K mutually exclusive states. Let $x_1, x_2 \dots x_k$ are indicator variables where x_i is 1 if the outcome is state i and 0 otherwise.
- **Gaussian distribution** – distributed with value of mean and variance.

CONTINUE...

- Assume, we are having **three data points** that has been generated using some data process. **Say, 9, 9.5 and 11.**
- Ultimately we are trying to calculate the total probability of observing all the data.
- Simply, the joint probability distribution of all the observed data points.
- The PDF for observing single data point x , using gaussian distribution:

$$P(x; \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

Data point Parameters of model

CONTINUE...

$$P(x; \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x - \mu)^2}{2\sigma^2}\right)$$

$$P(9, 9.5, 11; \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(9 - \mu)^2}{2\sigma^2}\right) \times \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(9.5 - \mu)^2}{2\sigma^2}\right) \\ \times \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(11 - \mu)^2}{2\sigma^2}\right)$$

$$\ln(P(x; \mu, \sigma)) = \ln\left(\frac{1}{\sigma\sqrt{2\pi}}\right) - \frac{(9 - \mu)^2}{2\sigma^2} + \ln\left(\frac{1}{\sigma\sqrt{2\pi}}\right) - \frac{(9.5 - \mu)^2}{2\sigma^2} \\ + \ln\left(\frac{1}{\sigma\sqrt{2\pi}}\right) - \frac{(11 - \mu)^2}{2\sigma^2}$$

$$\ln(P(x; \mu, \sigma)) = -3 \ln(\sigma) - \frac{3}{2} \ln(2\pi) - \frac{1}{2\sigma^2} [(9 - \mu)^2 + (9.5 - \mu)^2 + (11 - \mu)^2]$$

CONTINUE...

$$\frac{\partial \ln(P(x; \mu, \sigma))}{\partial \mu} = \frac{1}{\sigma^2} [9 + 9.5 + 11 - 3\mu] .$$

$$\mu = \frac{9 + 9.5 + 11}{3} = 9.833$$

MODEL SELECTION AND GENERALIZATION

- To understand learning taking example of Boolean function. Where all the inputs and outputs are binary.
- There are 2^d possible ways to write d binary values and so with d inputs, the training set has at most 2^d examples.

x_1	x_2	h_1	h_2	h_3	h_4	h_5	h_6	h_7	h_8	h_9	h_{10}	h_{11}	h_{12}	h_{13}	h_{14}	h_{15}	h_{16}
0	0	0	0	0	0	0	0	0	0	1	1	1	1	1	1	1	1
0	1	0	0	0	0	1	1	1	1	0	0	0	0	1	1	1	1
1	0	0	0	1	1	0	0	1	1	0	0	1	1	0	0	1	1
1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1

With two inputs there are four possible cases and sixteen possible Boolean functions

CONTINUE...

- Here each training examples removes half of the hypotheses, specifically those whose guesses are wrong.
- For example, in the case of $x_1=0$ and $x_2=1$ the output is 0. This removes $h_5, h_6, h_7, h_8, h_{13}, h_{14}, h_{15}, h_{16}$.
- This is one of the ways to see and understand the learning. With more examples we removes those hypotheses that are not consistent with the training data.
- In our cases of Boolean function, to end us with single hypothesis we need to see all the 2^d training examples.
- If training set contains only a small subset of possible instances, means the known true output has very small set of case to justify itself, then the solution is not unique.

CONTINUE...

- After seeing all the N examples cases, there remain 2^{2d-N} possible functions.
- This is referred as **ill-posed problem** – Where the data by itself are not sufficient to find a unique solution.
- The same problem also exists in many learning applications, in classification and in regression as well.
- So the learning is ill-posed the data itself are not sufficient to find the solution, we required to make some additional assumptions to have a specific unique solution with the existing data only, that we are having.
- The set of assumptions we have to make learning possible is called **inductive bias** of learning algorithm.

CONTINUE...

- Thus learning is not possible without inductive bias, and the process to choose right bias is called **model selection**.
- We need to understand that the scope of the machine learning is rarely to replicate the training data but the prediction for new cases.
- How well the model trained on the training set predicts the right outcome for new instances is called **generalization**.
- For the best generalization we should match the complexity of the hypothesis with the complexity of the function that underlying the data.

CONTINUE..

- If the hypothesis is less complex than the function, we have **underfitting**. As we try to fit a line on data sampled from third order polynomial.
- In such cases, as we increase the complexity both the training error and validation error decrease.
- But, if we have a hypothesis that is too complex, the data is not enough to constraint it and we may end up with a bad hypothesis.
- For example, fitting sixth order polynomial to noisy data sampled from a third order polynomial. This is called **overfitting**. In such case, having more training data helps but only up to a certain point.

CONTINUE...

➤ **The triple trade-off** – All learning algorithms those are trained from example data, there is a trade-off between three factors:

1. The complexity of hypothesis we fit to data, means the capacity of the hypothesis class.
2. The amount of training data.
3. The generalization error on new examples.

CONTINUE...

- The amount of training data increases, the generalization error decreases.
- The complexity of the model increases, the generalization error decreases first and then start to increase.
- The generalization error of a complex hypothesis can be kept in check by increasing the amount of training data but only up to a point.
- We can measure the generalization ability of a hypothesis, namely the quality of its inductive bias, if we have access to data outside the training set.
- We perform this by dividing the **training set** into two parts, we use one part for training and other called **validation set**, which is used to test generalization ability.

CONTINUE...

- Considering the large enough training and validation set, the hypothesis that is the most accurate on the validation set is considered as best one – the one that has the best inductive bias. The process is called **cross-validation**.
- In the case, when we need to report the error to give an idea about the expected error of our best model, we should not use the validation error.
- As we have used the validation set to choose the best model, and it has effectively become a part of training set.
- We need a third set, which should not be a part of training or validation set – called **test set**.

EVALUATING AN ESTIMATOR WITH BIAS AND VARIANCE

- Let X be a sample from a population specified up to a parameter θ and let $d=d(x)$ be an estimator of θ .
- To evaluate, the quality of this estimator, we can measure how much it is different from θ , that is $(d(x) - \theta)^2$. As it is a random variable it depends of the samples, we need to average it over possible X and consider $r(d, \theta)$ and so the mean square error of an estimator d defined as:


$$r(d, \theta) = E[(d(x) - \theta)^2]$$

The bias of an estimator is given as,

$$b_{\theta}(d) = E[d(x)] - \theta$$

CONTINUE...

- θ is the parameter to be estimated.
- d_i are several estimates over different samples.
- Bias is the difference between the expected value of d and θ .
- Variance is how much d_i are scattered around the expected value.
- **We would like both to be small.**





REFERENCES

[1] Introduction to Machine Learning – Ethem Alpaydin, MIT Press